


ECEN 4532 - Lab 4: MPEG Audio Signal Processing (MP3)

Andrew Teta

March 18, 2019



Contents

Introduction	3
1 Introduction	3
2 Background	3
2.1 Analysis	3
2.2 Synthesis	3
3 Cosine Modulated Pseudo Quadrature Mirror Filter: Analysis	4
3.1 The math	4
3.2 Derivation of a faster way	5
3.3 Frequency inversion	6
3.4 Implementation	7
4 Reconstruction: Synthesis	13
4.1 Calculations	13
4.2 Implementation	13
4.3 Selective sub-bands	18
5 Conclusion	18
6 Appendix	19
6.1 PQMF	19
6.2 IPQMF	19
6.3 Main	20

1 Introduction

In this lab, we explore signal processing in the context of audio compression. Specifically, we will be constructing the basis of the MPEG 1 Layer III codec, commonly known as MP3. The general idea is to implement a series of sub-band filters to decompose and reconstruct input audio "perfectly" and later, experiment with the exclusion of certain frequency bands as a method of data compression. We will implement the polyphase pseudo-QMF filter bank, which filters the audio signal into 32 frequency bands. The two main procedures in this processing scheme are analysis and synthesis.

2 Background

2.1 Analysis

Analysis is the process of decomposing a signal into frequency components and filtering it into sub-bands. MP3 encoding splits the audio signal $x(n)$ into 32 equally spaced frequency bands. This is accomplished with 32 parallel filters with impulse response h_k such that,

$$s'_k(n) = h_k(n) * x(n), \quad k = 1, \dots, 32. \quad (1)$$

By downsampling the output of each filter by a factor of 32, we achieve a critically sampled analysis filter. Thus,

$$s_k(n) = ((h_k * x) \downarrow 32)(n) = s'_k(32n). \quad (2)$$

The effect of this procedure (and the idea of critical sampling) is an efficient filter, where for an input block of size N samples of $x(n)$, a total of N output samples are produced.

2.2 Synthesis

Reconstruction is performed by upsampling the downsampled signals with zero values and filtering with a synthesis filter bank, g_k .

$$s_k \uparrow 32) * g_k(n), \quad k = 1, \dots, 32. \quad (3)$$

The synthesized output, \tilde{x} is found by the sum of all the synthesis filters,

$$\tilde{x} = \sum_{k=1}^{32} (s_k \uparrow 32) * g_k(n). \quad (4)$$

The MP3 filter bank has nearly (but not exactly) perfect reconstruction and fortunately the error is not an issue. The filter design is greatly simplified by using a prototype filter, modified only slightly for each of the 32 analysis filters, h_k .

3 Cosine Modulated Pseudo Quadrature Mirror Filter: Analysis

This section describes the mathematical description of MP3 analysis and further derives a fast algorithm to compute both convolution and decimation together, producing the output signal s_k for each of the 32 sub-bands.

3.1 The math

Consider a 512 tap filter h_k such that

$$s_k(n) = \sum_{m=0}^{511} h_k(m)x(32n - m) \quad (5)$$

and

$$h_k(m) = p_0(m) \cos \left(\frac{(2k+1)(m-16)\pi}{64} \right) \quad k = 0, \dots, 31, \quad m = 0, \dots, 511. \quad (6)$$

p_0 is a prototype lowpass filter (see Fig. 1) and the multiplication with $\cos \left(\frac{(2k+1)(m-16)\pi}{64} \right)$ shifts the center frequency of p_0 to be centered around $(2k+1)\pi/64$. Thus, h_k becomes a bandpass filter, selecting frequencies around $(2k+1)\pi/64$, for $k = 0, \dots, 31$, and a nominal bandwidth of $\pi/32$.




Figure 1: A cartoon representation of p_0 , the prototype filter used for analysis of a signal in MP3 encoding. This plot has no relevance to scale and is only a reference to the general shape of p_0 and its lowpass filter construction.

3.2 Derivation of a faster way

We begin with (5) and decompose the summation by letting $m = 64q + r$

$$s_k(n) = \sum_{m=0}^{511} h_k(m)x(32n - m) = \sum_{q=0}^7 \sum_{r=0}^{63} h_k(64q + r)x(32n - 64q - r) \quad (7)$$

and since

$$h_k(m) = p_0(m)\cos\left(\frac{(2k+1)(m-16)\pi}{64}\right), \quad (8)$$

we can write

$$h_k(64q + r) = p_0(64q + r)\cos\left(\frac{(2k+1)(64q + r - 16)\pi}{64}\right) \quad (9)$$

$$= p_0(64q + r)\cos\left(\frac{(2k+1)(r-16)\pi}{64} + (2k+1)q\pi\right) \quad (10)$$

$$= \begin{cases} p_0(64q + r)\cos\left(\frac{(2k+1)(r-16)\pi}{64}\right) & \text{if } q \text{ is even} \\ -p_0(64q + r)\cos\left(\frac{(2k+1)(r-16)\pi}{64}\right) & \text{if } q \text{ is odd} \end{cases} \quad (11)$$




Figure 2: Plot of a 512 tap analysis filter, C , used in MP3 encoding.

Let us define (see Fig. 2),

$$c(m) = \begin{cases} p_0(m) & \text{if } q = \lfloor m/64 \rfloor \text{ is even} \\ -p_0(m) & \text{if } q = \lfloor m/64 \rfloor \text{ is odd,} \end{cases} \quad (12)$$

then

$$h_k(64q + r) = c(64q + r)\cos\left(\frac{(2k+1)(r-16)\pi}{64}\right). \quad (13)$$

Using the notation of the MP3 standard, we define

$$M_{k,r} = \cos\left(\frac{(2k+1)(r-16)\pi}{64}\right), \quad k = 0, \dots, 31, \quad r = 0, \dots, 64, \quad (14)$$

so,

$$h_k(64q + r) = c(64q + r)M_{k,r}, \quad (15)$$

and

$$s_k(n) = \sum_{q=0}^7 \sum_{r=0}^{63} M_{k,r} c(64q + r) x(32n - 64q - r) \quad (16)$$

Expressing the analysis in the form of (16), we can efficiently compute sub-band sample coefficients. For every n , we compute:

$$z(64q + r) = c(64q + r) x(32n - 64q - r), \quad r = 0, \dots, 63, \quad q = 0, \dots, 7, \quad (17)$$

sum over q

$$y(r) = \sum_{q=0}^7 z(64q + r), \quad r = 0, \dots, 63, \quad (18)$$

and compute one sample output for each sub-band by taking a sum over r

$$s_k = \sum_{r=0}^{63} M_{k,r} y(r), \quad k = 0, \dots, 31. \quad (19)$$

3.3 Frequency inversion

Downsampling has the effect of moving the output of each filter down to the baseband. However, the periodic nature of the discrete-time Fourier transform (DTFT) representation, causes a "frequency inversion" to occur, where higher frequency components appear to be to the left of the y-axis. This can be illustrated visually in Fig. 3.

Considering the first sub-band filter to be index 0, then as it works out, the *odd* sub-bands are inverted. The MP3 codec further decomposes the filter bank output using the modified discrete cosine transform (MDCT), so it is desirable to undo the inversion. In complex exponential form, this can be accomplished by multiplying the signal by $e^{j\pi n}$, but in Python and other analysis tools where only the exponential coefficients are considered, we can simply multiply the sample output of every odd sub-band by -1. This operation has the effect of translating the spectral signal by π , restoring its orientation because the DTFT is periodic in 2π .




Figure 3: Spectrum inversion

3.4 Implementation

Our analysis was done in Python, however the basic procedure of implementation will be the same for any implementation.

Each processing cycle works on packets of 32 audio samples and the filtering is performed on a buffer X of size 512. The buffer is shifted to the right by 32 samples and the available indexes at the left end of the buffer are filled with new audio samples. 32 sub-band coefficients are computed and this process is repeated until the length of the input signal is reached.

In this lab, the window filter coefficients, c , were calculated prior and provided in the form of a text file. The 512 coefficients were read into a `numpy` array at the beginning of analysis to be used in calculation later. Audio was provided as `.wav` tracks. We used `scipy.io.wavfile.read()` to convert these files to `numpy` arrays. Further analysis was performed on the first 5 seconds of each track.

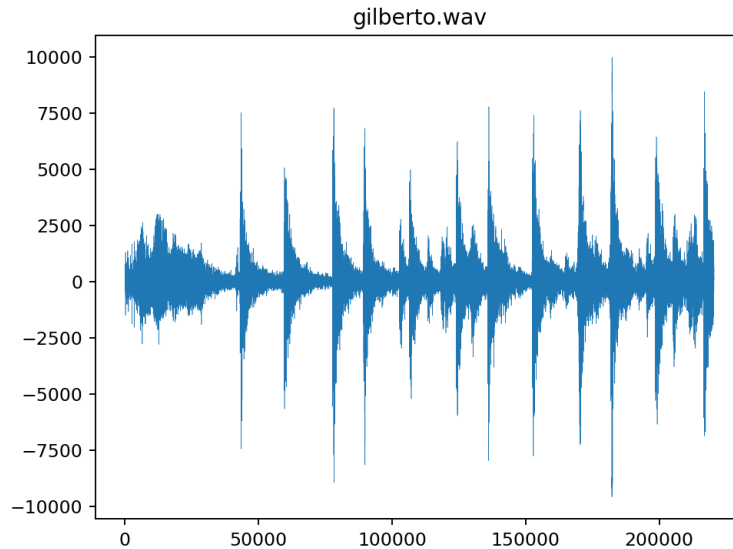


Figure 4: Plot of samples vs. magnitude for first 5 seconds of a jazz song, used as input to our analysis.

We begin calculation of sub-band coefficients by reshaping the audio sample array to have dimension $(-1,32)$, where -1 means the first dimension is implied. I implemented a quick calculation to append zeros to the end of the array if it did not have the correct number of elements to be reshaped directly. Before beginning to loop over the rows of this new matrix of audio samples, we can build M one time to save processing cycles. A temporary buffer, X of length 512, an output matrix of `np.zeros_like` the $(-1,32)$ input sample data, and a couple helper vectors which allow vector operations instead of sums, are all declared outside of the loop. One helper vector, `fInvert` is used for frequency inversion correction.

$$fInvert = [1, -1, 1, \dots, 1] \quad \text{length} = 32$$

As well as `zFlat`, a length 8 vector of ones which replaces the sum in (18).

Now for the actual analysis. We loop over each row, effectively performing operations on a "packet" of 32 samples for each loop iteration. For each packet,

1. Shift every value in buffer X right by 32
2. Fill $X[0:32]$ with packet (flipped to perform convolution)
3. Compute $Z = C * X$, windowing X by the analysis filter taps (eq. (17))
4. Reshape Z into $(8,64)$ and calculate $Y = zFlat \cdot Z$ (eq. (18))
5. Compute $S = M \cdot Y$ (eq. (19))
6. Store length 32 output vector, S in a row of output matrix

The output of this sequence should be a matrix of sub-band coefficients organized like so (input matrix of shape $(N,32)$):

$$A = \begin{bmatrix} sb_{0,0} & sb_{0,1} & sb_{0,2} & \dots & sb_{0,31} \\ sb_{1,0} & sb_{1,1} & sb_{1,2} & \dots & sb_{1,31} \\ \dots & & & & \\ sb_{N,0} & sb_{N,1} & sb_{N,2} & \dots & sb_{N,31} \end{bmatrix} \quad (20)$$

and will also have shape $(N,32)$.

Transposing and flattening A , we can plot the sub-band analysis output.




Figure 5: Spectral content of a piano melody. We see a small magnitude of very low frequencies, most of the energy concentrated in the low-mid frequency range, and almost no content at high frequency.




Figure 6: Spectral content of a dance floor beat. We see high energy density at very low frequencies and lots of noise at higher frequencies.

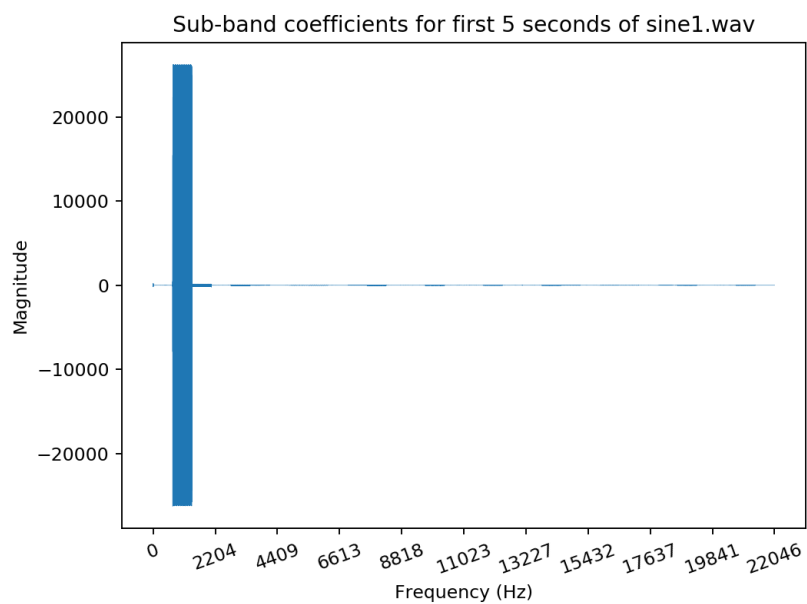


Figure 7: Spectral content of a sine wave. We see energy only between about 750 Hz and 1250 Hz.




Figure 8: Spectral content of a sine wave. We see energy only between 0 Hz and 750 Hz.

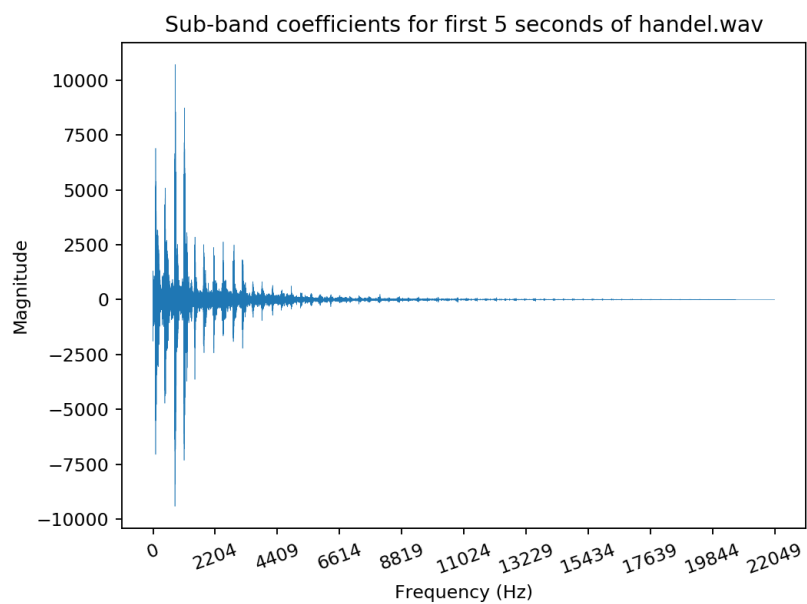


Figure 9: Spectral content of a classical track. There is good contribution from spectral bands at 0 Hz all the way up to about 5 kHz.

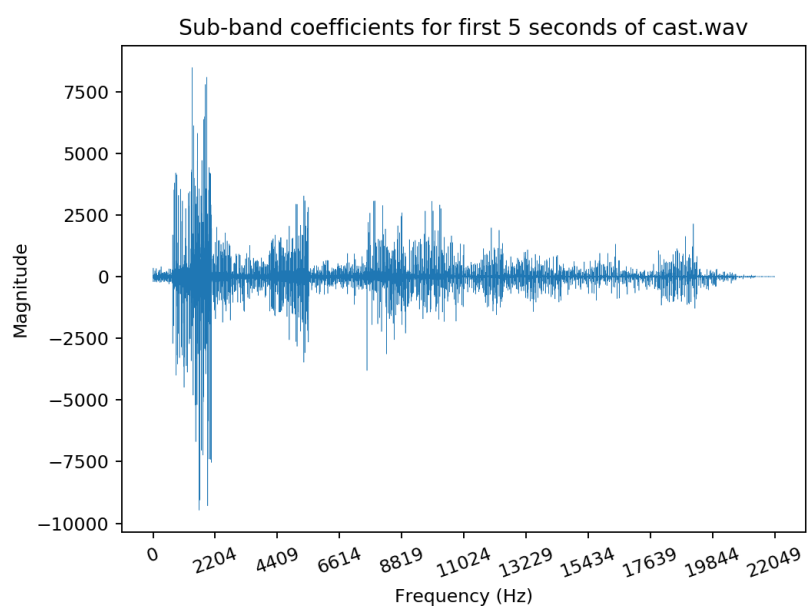


Figure 10: Spectral content of a percussive recording. The energy is distributed fairly evenly across the frequency spectrum here.

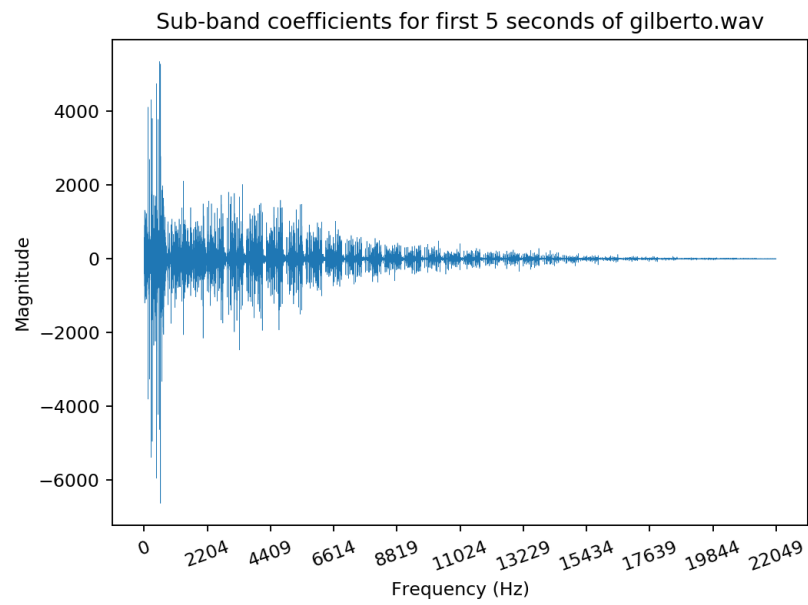


Figure 11: Spectral content of a jazz track. The first five seconds are very percussion heavy, most likely leading to the large spread of spectral content. Above about 16 kHz there is mostly noise.

4 Reconstruction: Synthesis

Synthesis, or reconstruction of the MP3 codec is very similar to analysis. For each set of 32 sub-band coefficients, 32 audio samples are reconstructed. The algorithm will need to reverse the frequency inversion correction, so we will actually invert the signal again.

4.1 Calculations

This time, we will perform modulation with a matrix N where,

$$N_{i,k} = \cos\left(\frac{(2k+1)(16+i)\pi}{64}\right), \quad i = 0, \dots, 63, \quad k = 0, \dots, 31. \quad (21)$$

First, we load reconstruction filter tap coefficients from a text file to obtain the window in Fig. 12




Figure 12: Plot of reconstruction window D .

We again operate on "packets" of length 32 samples and loop over vectors of this length through to the end of the track. This time our buffer vector V is of length 1024, we declare length 512 vectors U and W , an output vector S and two more helping vectors, `fInvert` as before and a length 16 vector of ones, `wFlat`.

We will build U during each cycle and use it to compute the reconstruction. V is constructed over multiple packet cycles and so the content of U is also updated. We define

```
for i = 0 to 7 do
  for j = 0 to 31 do
    u[64 i + j]      = v[128 i + j]
    u[64 i + j + 32] = v[128 i + j + 96]
```

4.2 Implementation

The steps in reconstruction go as follows:

1. Shift every value in a buffer V right by 64
2. Calculate modulation and re-invert: $V[0 : 64] = N \cdot (fInvert * packet)$
3. Compute U as shown above
4. Window the samples vector by $W = U * D$

5. Reshape W into $(16,-1)$ and store $wFlat \cdot W$ in output matrix

Plotting the output of this reconstruction procedure, we can compare it to the original signal.




Figure 13: Plot of first 1000 samples of a simple sine wave input signal, `sine1.wav`. Observe the reconstructed signal is delayed by some amount.

error = 2.62

Through trial and error, attempting to minimize the error between the original signal and reconstruction, I found the delay to be equal to 481 samples. This is caused by having an empty buffer when we begin synthesis. The it is only after 15 packets of information, that the buffer is full (with 512 values) and so the first 15 cycles produce a delay of available information. We start with 32 samples and index from 0, thus we have a delay of 481 samples.




Figure 14: Reconstruction and delay shown for sample1.wav
error = 0.24




Figure 15: Reconstruction and delay shown for sample2.wav
error = 1.32




Figure 16: Reconstruction and delay shown for sine2.wav
error = 1.47




Figure 17: Reconstruction and delay shown for handel.wav
error = 0.15




Figure 18: Reconstruction and delay shown for cast.wav
error = 0.0




Figure 19: Reconstruction and delay shown for gilberto.wav
error = 0.065

The errors listed in each of the figures above are calculated as maximum error between the original samples and reconstructed samples for only the first 1000 samples. It is scaled based on the maximum sample value in each song, so the values appear much larger than they are if the data were normalized.

4.3 Selective sub-bands

Since MP3 is a compression algorithm, we are interested in reconstructing partial data less than perfectly. In the MP3 codec, this is implemented by applying a psychoacoustic model to select a subset of the 32 sub-bands based on human auditory sensitivity. For this lab, we will manually select some bands to leave out, attempting to limit the data used to reconstruct without negatively affecting the perception of music.

For some tracks this is trivial, while for others it is a more complicated task. We will define a compression to be

$$\frac{\text{number of bands used to reconstruct}}{32} \quad (22)$$

Selecting bands only where the frequency has large magnitude, the error was calculated as follows for each track:

$$cast = 0.83$$

$$gilberto = 229.88$$

$$handel = 164.11$$

$$sample1 = 115.99$$

$$sample2 = 10376.86$$

$$sine1 = 789.02$$

$$sine2 = 6735.96$$

Again, these errors are not scaled to normalization and a few are probably much larger than ideal, but this was not very analytic.

5 Conclusion

This lab was conceptually much more difficult than the previous labs and the code was really nothing new. I did learn a lot more about the DFT, DTFT, frequency analysis in general, and how all of that can be implemented in software. I feel like my foundation of Fourier analysis is much stronger now. This was a cool introduction to a real compression scheme.

6 Appendix

6.1 PQMF

```
1 def pqmf(input):
2
3     # load window coefficients from file
4     c_taps = load_c_taps('the_c_taps.txt')
5
6     plt.figure(dpi=170)
7     plt.plot(c_taps)
8     plt.savefig('./figures/analysis/c_taps.png')
9     plt.close()
10
11     # add zeros at end of array to make it reshappable into 32 cols
12     data = np.append(input, np.zeros(32 - (len(input) % 32)))
13     # reshape into 32 sample length rows
14     data = np.reshape(data, (-1, 32))
15
16     # initialize working arrays for subband coefficient calculation
17     X = np.zeros(512)
18
19     # build filter matrix, M
20     M = np.zeros([32, 64])
21     for k in range(32):
22         for r in range(64):
23             M[k, r] = math.cos(((2 * k) + 1) * (r - 16) * math.pi) / 64)
24
25     # subband coefficient output matrix
26     A = np.zeros_like(data)
27
28     # helping vectors
29     fInvert = np.array(16 * [1, -1])
30     zFlat = np.array(8 * [1])
31
32     # loop over entire song in 32 sample chunks
33     for packet in range(np.shape(data)[0]):
34         # shift every element of X to the right by 32
35         X = np.roll(X, 32)
36         # flip audio sample frame and place in X
37         X[0:32] = np.flip(data[packet])
38         # window X by C filter
39         Z = c_taps * X
40         # partial calculation
41         Z = np.reshape(Z, (8, 64))
42         Y = zFlat.dot(Z)
43         # calculate 32 subband samples
44         S = M.dot(Y)
45         # undo frequency inversion and add to output array
46         A[packet, :] = fInvert * S
47
48     return A
```

6.2 IPQMF

```
1 def ipqmf(input, subbands):
2
3     data = input
4
5     # load window coefficients from file
6     d_taps = load_d_taps('the_d_taps.txt')
7
8     plt.figure(dpi=170)
9     plt.plot(d_taps)
10    plt.savefig('./figures/synthesis/d_taps.png')
11    plt.close()
12
13    # declare working array
14    V = np.zeros(1024)
15    U = np.zeros(512)
16    W = np.zeros(512)
17
18    # reconstruction coefficient output matrix
19    S = np.zeros_like(data)
20
21    # build reconstruction matrix, N
22    N = np.zeros([64, 32])
23    for i in range(64):
24        for k in range(32):
25            N[i,k] = math.cos(((2 * k) + 1) * (16 + i) * math.pi) / 64)
26
27    # helping vectors
28    fInvert = np.array(16 * [1, -1])
29    wFlat = np.array(16 * [1])
30
31    # loop over coefficients in 32 sample chunks
32    for packet in range(np.shape(data)[0]):
33        # filter output sub-bands
34        data[packet] = data[packet]
35        # shift every element of V to the right by 64
36        V = np.roll(V, 64)
37        # compute reconstruction samples
38        V[0:64] = N.dot(fInvert * data[packet] * subbands)
39        # build window operand
40        for i in range(8):
41            for j in range(32):
42                U[i * 64 + j] = V[i * 128 + j]
43                U[i * 64 + 32 + j] = V[i * 128 + 96 + j]
44        # window
45        W = U * d_taps
46        W = np.reshape(W, (16, -1))
47        S[packet, :] = wFlat.dot(W)
48
49    return S
```

6.3 Main

```
1 import lab04_funcs as lf
2 from tkinter import filedialog
3 import numpy as np
4 import scipy.io.wavfile
5 import matplotlib.pyplot as plt
6
7 d_taps = lf.load_d_taps('the_d_taps.txt')
8
9 # UI dialog to select files -> selection of multiple files will run all functions for
  each file
10 files = filedialog.askopenfilenames()
11 # Loop over files selected
12 for f in files:
13     filename = str.split(f, '/')[−1]
14     filename = str.split(filename, '.')[0]
15     filepath = f
16
17     # read .wav file into numpy array and extract sampling frequency
18     fs, data = scipy.io.wavfile.read(filepath)
19
20     # extract first 5 seconds of .wav file
21     data = data[0 : int(5 * fs)]
22
23     # plot input samples
24     plt.figure(dpi=170)
25     plt.plot(data, linewidth=0.25)
26     plt.title(filename + '.wav')
27     plt.savefig('./figures/' + filename + '.png')
28     plt.close()
29
30     # calculate sub-band filter coefficients
31     print('Analyzing ' + filename + ' ... \n')
32     coefficients = lf.pqmf(data)
33
34     # vector to filter subbands
35     thebands = np.ones(32)
36
37     # calculate reconstruction coefficients
38     print('Reconstructing ' + filename + ' ... \n')
39     recons = lf.ipqmf(coefficients, thebands)
40
41     # transpose and flatten to sort data into incrementing groups of subband
  coefficients
42     coefficients = coefficients.T.flatten()
43
44     # plot sub-band coefficients
45     plt.figure(dpi=170)
46     plt.plot(coefficients, linewidth=0.25)
47     xlocs = np.asarray(range(0, len(coefficients), int(len(coefficients)/10)))
48     xvals = np.asarray(xlocs/np.amax(range(len(coefficients)))*fs/2, int)
49     plt.xticks(xlocs, xvals, rotation=20)
50     plt.xlabel('Frequency (Hz)')
51     plt.ylabel('Magnitude')
52     plt.title('Sub-band coefficients for first 5 seconds of ' + filename + '.wav')
53     plt.tight_layout()
54     plt.savefig('./figures/analysis/' + filename + '.png')
55     plt.close()
56
57     recons = recons.flatten()
58
59     # plot reconstruction coefficients
60     plt.figure(dpi=170)
61     plt.plot(recons, linewidth=0.25)
62     plt.xlabel('Sample')
```

```

63 plt.ylabel('Magnitude')
64 plt.title('Reconstruction of first 5 seconds of ' + filename + '.wav')
65 #plt.show()
66 plt.savefig('./figures/synthesis/' + filename + '.png')
67 plt.close()
68
69 delay = 512 - 31;
70
71 # calculate error
72 r = recons[delay:1000 + delay]
73 d = data[0:1000]
74 error = np.amax(r - d)
75 print(filename + f' error = {error}')
76
77 # plot delay
78 plt.figure(dpi=170)
79 plt.plot(data[0:1000], linewidth=0.35)
80 plt.plot(recons[0:1000], linewidth=0.35)
81 if (filename == 'cast'):
82     plt.clf()
83     plt.plot(data[2850:3850], linewidth=0.35)
84     plt.plot(recons[2850:3850], linewidth=0.35)
85 plt.title('Original waveform vs. reconstructed')
86 plt.savefig('./figures/synthesis/' + filename + '_delay.png')
87 plt.close()
88
89 thebands = np.ones(32)
90 coefficients = np.reshape(coefficients, (32, -1)).T
91 if filename == 'gilberto':
92     thebands[16:32] = 0
93 elif filename == 'sinel':
94     thebands[0] = 0
95     thebands[2:32] = 0
96 elif filename == 'sine2':
97     thebands[1:32] = 0
98 elif filename == 'handel':
99     thebands[8:32] = 0
100 elif filename == 'sample1':
101     thebands[8:32] = 0
102 elif filename == 'sample2':
103     thebands[15:32] = 0
104 selective_output = lf.ipqmf(coefficients, thebands)
105 selective_output = selective_output.flatten()
106 # calculate error
107 r = selective_output[delay:delay + 5000]
108 d = data[0:5000]
109 error = np.amax(r - d)
110 print(filename + f' selective error = {error}')
111
112
113 print('done')

```