

The Dot Product

CU Boulder

- ▶ The “dot product” is a key concept! It is central to geometric interpretations of matrix operations. Two vectors are orthogonal if their dot product is zero

$$\underline{\mathbf{x}}_1 \cdot \underline{\mathbf{x}}_2 = \sum_{i=1}^n \underline{\mathbf{x}}_1(i) \underline{\mathbf{x}}_2(i)$$

- ▶ We also have:


$$|\underline{\mathbf{x}}| = \sqrt{\underline{\mathbf{x}} \cdot \underline{\mathbf{x}}}$$

$$a\underline{\mathbf{x}} \cdot b\underline{\mathbf{y}} = (ab)(\underline{\mathbf{x}} \cdot \underline{\mathbf{y}}) = ab\underline{\mathbf{x}} \cdot \underline{\mathbf{y}}$$

Dot Product and Projection

CU Boulder

- ▶ If \underline{u} is a unit vector, the dot product $\underline{x} \cdot \underline{u}$ tells us the projection of \underline{x} in the direction of \underline{u}



Orthonormal Basis

CU Boulder

- ▶ **A set of vectors forms an orthonormal basis if**
 - Each vector is a unit vector (magnitude 1)
 - Each vector is orthogonal to every other vector

Mathematically,

$$\underline{\mathbf{u}}_i \cdot \underline{\mathbf{u}}_j = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{else} \end{cases}$$

Transforms are Simple! Really!

CU Boulder

- ▶ A transform like the DFT (and there are many, many different transforms) is defined by a set of basis vectors
- ▶ The vectors are orthonormal
- ▶ Any vector can be projected onto each of these basis vectors to get the transformed values, i.e., transform coefficients
- ▶ NOTE: the dot product of \underline{x} with \underline{y} when \underline{x} and \underline{y} are complex valued is \underline{x} “dot” \underline{y}^* where \underline{y}^* is the complex conjugate vector (conjugate each component)

Matrix description of a Transform

CU Boulder

- ▶ **Transforms (“change of basis”) can be written as a matrix equation**

- Make the basis vectors the rows of a “transform matrix”

$$\underline{\mathbf{c}} = \underline{\mathbf{A}} \underline{\mathbf{x}}$$

The transform coefficients

$$= \begin{bmatrix} \underline{\mathbf{u}}_1^T \\ \underline{\mathbf{u}}_2^T \\ \vdots \\ \underline{\mathbf{u}}_N^T \end{bmatrix} \underline{\mathbf{x}}$$

Vector being “transformed” (i.e. for which we are computing projections onto a new basis set)

- Every set of orthonormal basis vectors forms a “transform” via a matrix multiply

Some DFT insights

CU Boulder

- ▶ **For the DFT, the basis vectors are complex sinusoids**
- ▶ **We use $N/2+1$ DFT coefficients for real-valued signals**
 - Complex conjugacy
- ▶ **The spectrogram (STFT)**
 - A series of spectral plots presented one after the other
 - ✓ A pure sinusoid appears as a straight line
 - Each spectral plot is one frame of data
 - The frames are “short” and may overlap

► Header contains

- Sampling frequency
 - ✓ 44.1 kHz, 22.05 kHz, 11.025 kHz, 8 kHz are common
- Number of channels
 - ✓ 1 and 2 channels are common
- Number of samples
 - ✓ A “sample” is all the data values for every channel
- Bits per audio data value
 - ✓ 16 bits and 8 bits are common
 - 16 bit signed integer format covers -32768 to 32767

► Reading and writing a wav file

```
[bombay% cat zuppy.py
#!/usr/local/bin/python3
import scipy.io.wavfile as spwav

#####
# Here is our main block of code.  Execution starts here
#####

Fs, b = spwav.read("test.wav")
print(f"the sampling rate is {Fs}")
print(f"the number of audio samples is {len(b)}")
print(f"b[50000] = {b[50000]}, b[100000] = {b[100000]}")
print(f"type( b ) = {type(b)}")
print(f"type( b[50000] ) = {type( b[50000] )}")
spwav.write("newtest.wav", Fs, b)
```

```
[bombay%
[bombay% zuppy.py
the sampling rate is 11025
the number of audio samples is 2317404
b[50000] = -2462, b[100000] = 1767
type( b ) = <class 'numpy.ndarray'>
type( b[50000] ) = <class 'numpy.int16'>
[bombay%
```

Code

Output

► Getting a list of files in a directory

- If you include “import glob” at the top of your python program then you can get a list of files matching some criteria with a call like this:
 - ✓ `needed_files = glob.glob("*.wav")`
- Extracting 24 seconds of audio
 - ✓ Compute the sample at which to start extracting via:
 - `num_samps // 2`
 - ✓ We want the number of samples extracted to be a multiple of 512 to simplify later analysis. If you extract 516 frames of size 512, this is nearly 24 seconds of data ($512 \times 516 / 11025 = 23.96$ s)

► Don't use a for loop to do the extraction!

- Learn to use python “slicing”. For example, if `x` is a 1-D numpy array, then `x[8:15]` is a new array with `x[8]` to `x[14]` inclusive
 - ✓ `x[:15]` gives all values up to `x[14]`
 - ✓ `x[25:]` gives all samples from `x[25]` to the end
- I recommend reshaping your 1-D array of audio samples into an array of size 516 x 512
 - ✓ Look at the command `numpy.reshape()`
 - ✓ If `x` is a re-shaped 2D array, then `x[n]` gives you the 512 samples of frame `n`
 - Note: this means python lets you drop the second index entirely when what you are after is `x[n,:]`

► Minimize for() loops!

- You can do all of your frame calculations for time domain and spectral features using a single for() loop that indexes frames
 - ✓ You don't need a nested for() loop that looks at individual samples
 - ✓ If you use nested for() loops you aren't using the power of numpy/python (your grade will reflect that)

► Slicing with -n

- `x[:-1]` gives an array with all samples in `x` except for the last sample. This will help with ZCR calculations
 - ✓ This syntax extends to `x[:-n]` when you want to drop `n` samples off the end

► Spectrogram calculation (IMPORTANT)

- Import this module
 - ✓ “import scipy.signal as spsig”
- Computing a window
 - ✓ `blackman_window = spsig.get_window("blackman", N, False)`
- Use “`spsig.spectrogram()`”
 - ✓ See the documentation on the web
 - ✓ You can specify the window, for example, ‘blackman’, when invoking the function
 - ✓ use: `mode=‘magnitude’`
- Do your plots using dB not linear!
 - ✓ See the sample code on the next slide

► Spectrogram calculation and graphing examples

- Make sure you know what `np.where()` is doing below. It is a super useful function for modifying arrays (also, since I used a 2D array, I "flattened" it before doing the call)

```
f, t, Sxx = spsig.spectrogram( tmpdat.flatten(), Fs, 'blackman', N, \
                               mode='magnitude')
```

```
Sxx = np.where(Sxx < 10**-3, 10**-3, Sxx)
Sxx = 20 * np.log10( Sxx / np.amax(Sxx) )
```

```
plt.figure(3)
# cmaps of interest: inferno, BuPu, coolwarm, bwr, Reds, gray, plasma
plt.pcolormesh(t, f, Sxx, cmap='coolwarm')
plt.ylabel("Frequency (Hz)" )
plt.xlabel("time (s)")
print(f"shape of Sxx is {Sxx.shape}")

plt.show()
```

► Generating multiple plots

- You can generate multiple plots in separate windows by inserting `"plt.figure(n)"` in front of the commands for graph `n`, and putting a single `plt.show()` command after all figures you want have been constructed
- You can generate multiple plots in a single window using `"plt.subplot()"`

More Hints

CU Boulder

- ▶ **Don't forget to apply the blackman filter to each frame of data when computing the spectral features**
- ▶ **ERROR IN WRITEUP:** the size of your array for the mel filter banks should be N_b by $K+1$ (the writeup says $N_b \times K$)
- ▶ **The frequency corresponding to the k 'th coefficient when doing a DFT is $k \cdot F_s / N$**
- ▶ **The graph in the lab writeup of the frequency responses is WRONG. It gives the right basic shape, however, the heights are wrong**